

Logo Detection Using PyTorch

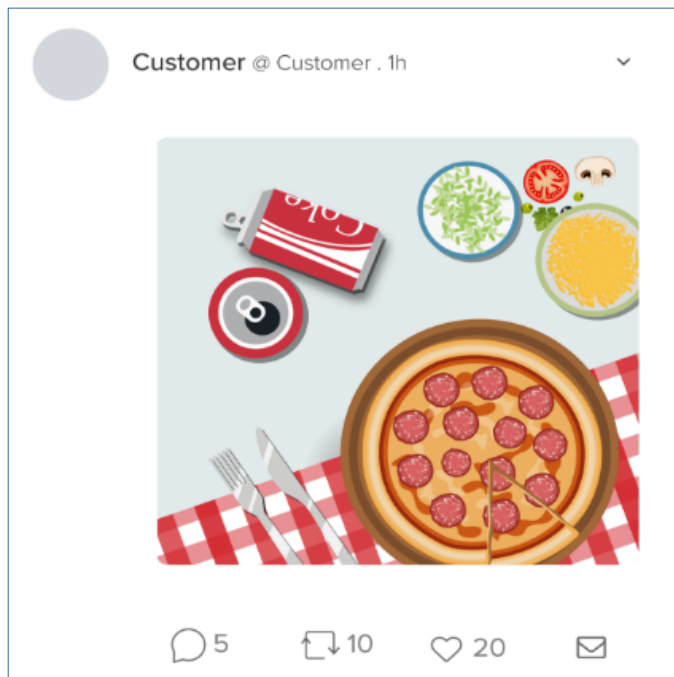
Nithiroj Tripatarasit (Lek)

Ad Tech

*Social
Medias*

Analyze

Target



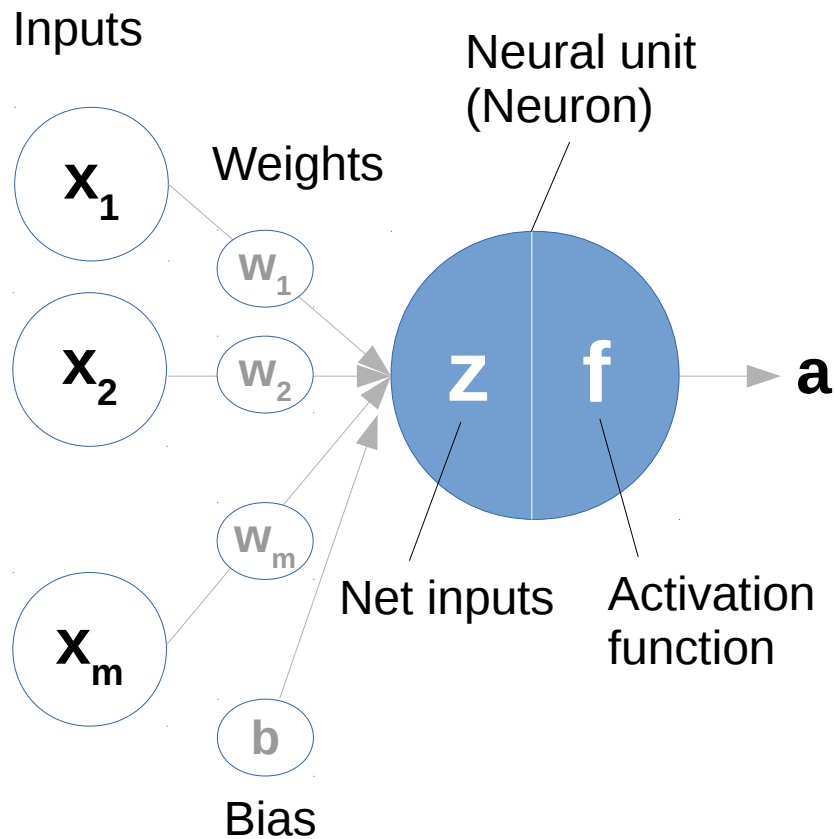
- Platform
- Category
- Brand
- Gender
- 18-24 years
- Age
- Location
- Ethnicity
- Weather



Deep Learning

- 1. Create the Network**
- 2. Train the Network**
- 3. Deploy the Network**

Single-Layer Neural Network




$$z = w_1 x_1 + w_2 x_2 + \dots + w_m x_m + b$$

$$f = \max(0, z)$$

ReLU

(Rectified Linear Unit)




$$a = f(z)$$

Fully Connected Network

Input layer with
 m input units

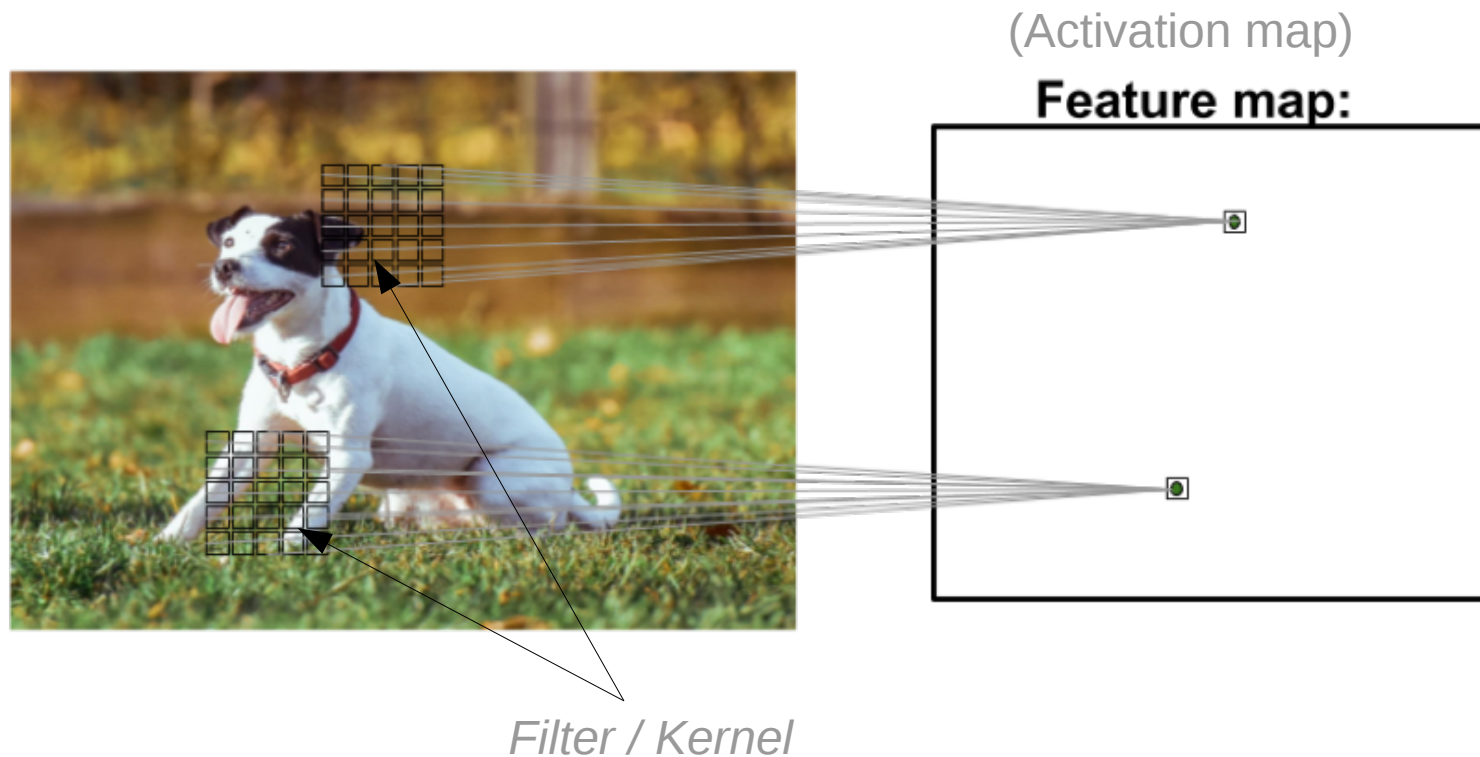
Hidden layer with
 n neural units

Output layer with
 k neural units




Convolutional Neural Networks (CNN)

It is an important tool for object recognition.



Convolutional Layer



Dot Product

INPUT
5 x 5 x 1

1	2	0	1	1
1	1	1	1	1
0	0	2	1	1
2	2	0	0	1
2	1	1	2	0

Filter
3 x 3 x 1
Stride = 1, Padding = 0

-1	1	0
-1	-1	0
1	0	0

weights

1

bias

OUTPUT
3 x 3 x 1
(Activation Map)

0	-3	2
3	1	-2
-1	2	1

$$(1 * -1) + (2 * 1) + (0 * 0) + (1 * -1) + (1 * -1) + (2 * 0) + (2 * 1) + (0 * 0) + ((1 * 0) + 1 = 0$$

Max Pooling

Single depth slice

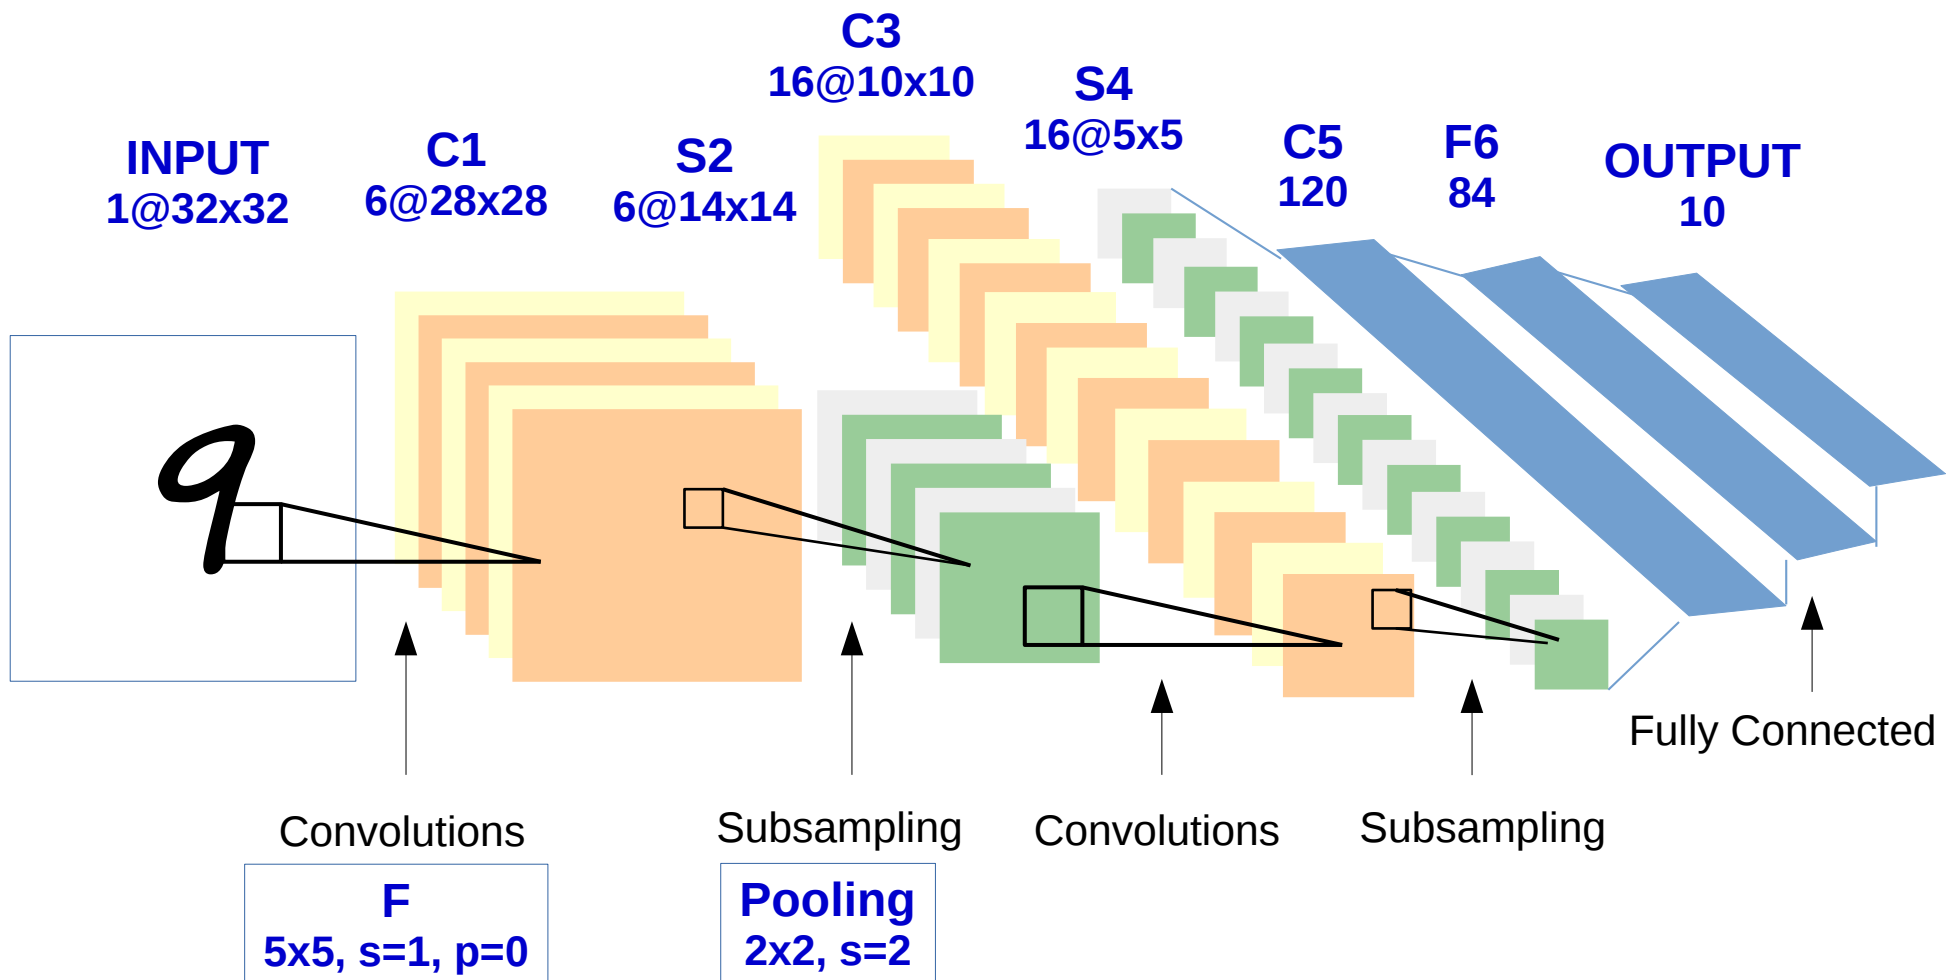
1	1	2	4
5	6	7	8
3	2	1	0
1	2	3	4

max pool with 2x2 filters
and stride 2

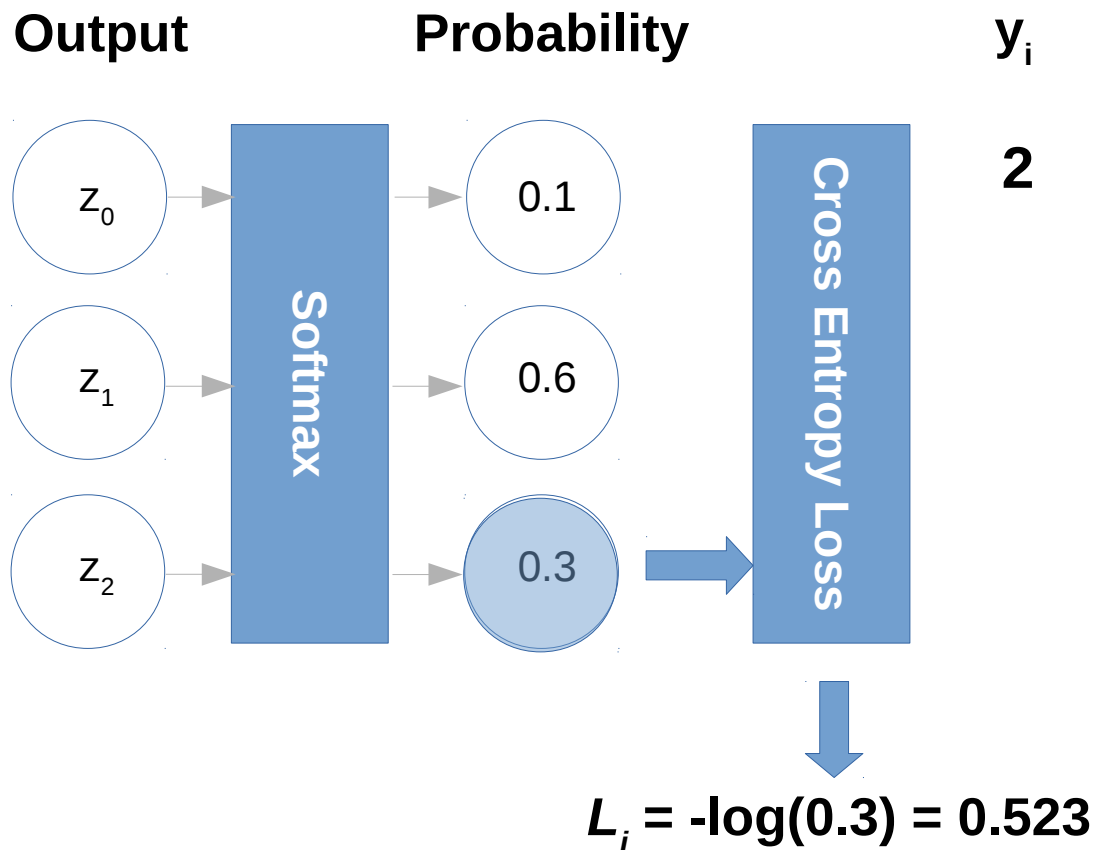


6	8
3	4

LeNet-5



Loss Function



Softmax


$$P(Y = k|X = x_i) = \frac{e^{s_k}}{\sum_j e^{s_j}}$$

Cross Entropy Loss

$$L_i = -\log\left(\frac{e^{sy_i}}{\sum_j e^{s_j}}\right)$$

Gradient Descent


An algorithm used to optimize the network



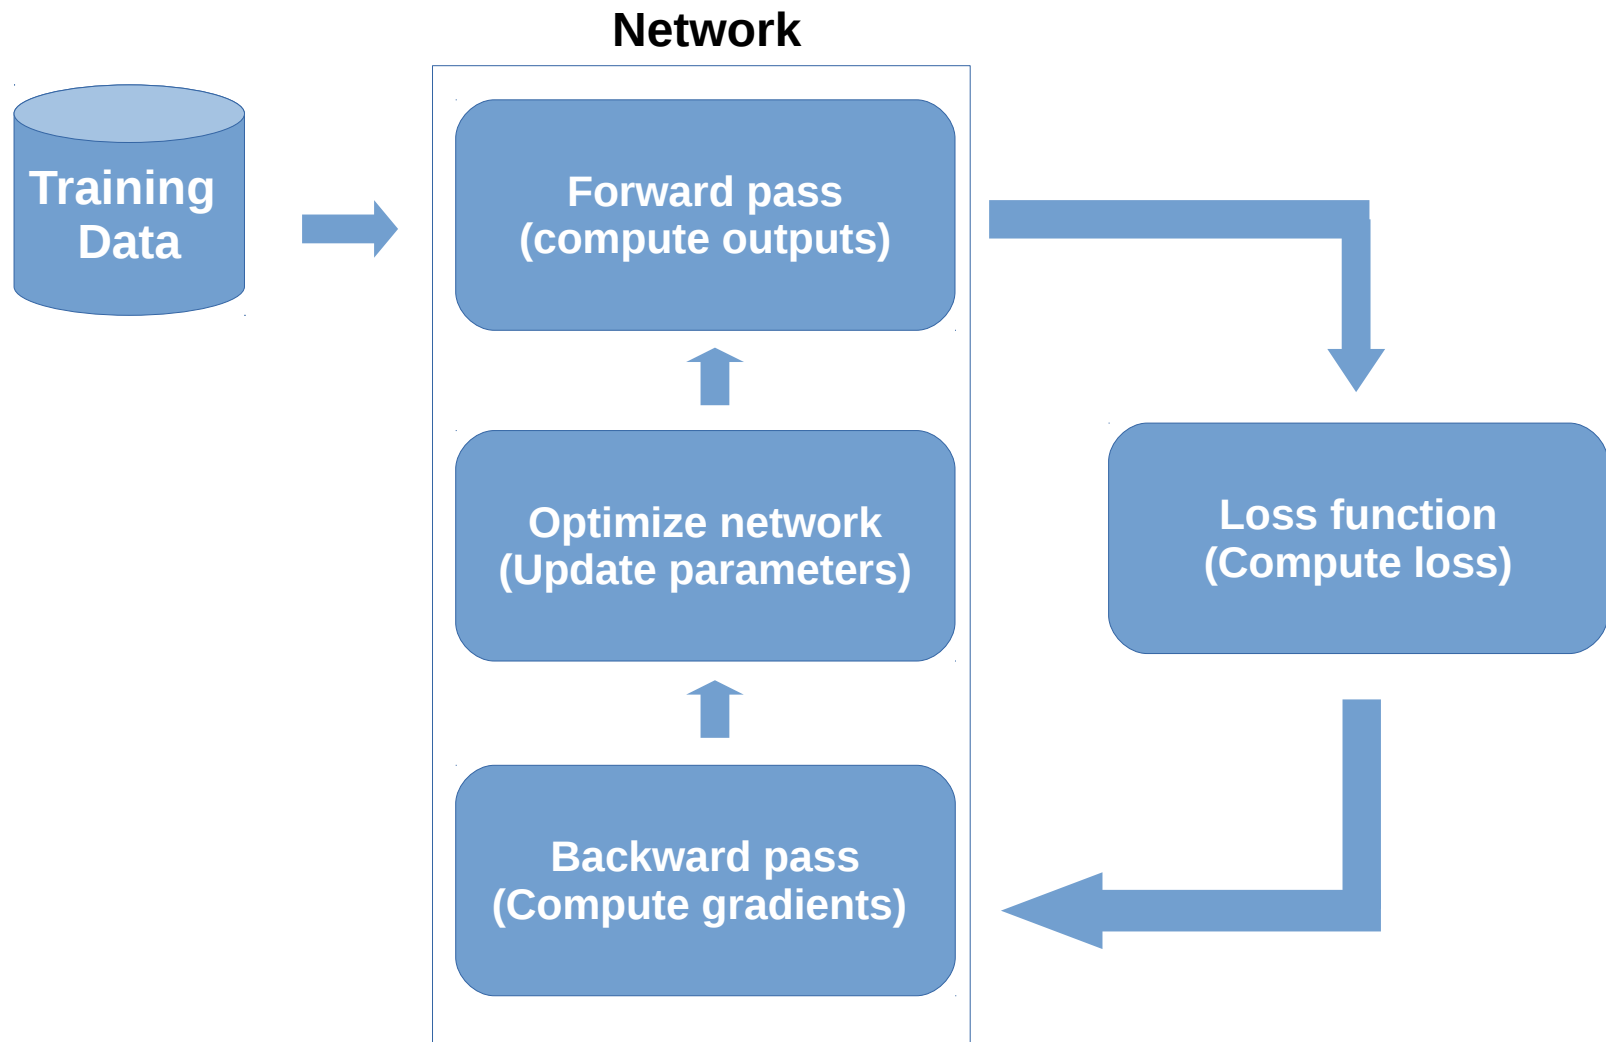
Update weights : $w += - \text{learning_rate} * \text{gradient}$

Optimizers

Algorithms used to optimize the network



Network Training Loop



Deep Learning Framework

PyTorch

- ◆ **Deep learning framework**
- ◆ **Autograd**
- ◆ **Tools to create and train deep learning easily and efficiently**
- ◆ **GPU support**

pytorch.org

Get Started.

Select your preferences, then run the PyTorch install command.

Please ensure that you are on the latest pip and numpy packages.
Anaconda is our recommended package manager

OS	<input checked="" type="radio"/> Linux	<input type="radio"/> MacOS	<input type="radio"/> Windows	
Package Manager	<input checked="" type="radio"/> conda	<input type="radio"/> pip	<input type="radio"/> Source	
Python	<input type="radio"/> 2.7	<input checked="" type="radio"/> 3.5	<input type="radio"/> 3.6	
CUDA	<input checked="" type="radio"/> 8	<input type="radio"/> 9.0	<input type="radio"/> 9.1	<input type="radio"/> None

Run this command:

```
conda install pytorch torchvision -c pytorch
```

[Click here for previous versions of PyTorch](#)

Tensors

't'
'e'
'n'
's'
'o'
'r'

tensor of dimensions [6]
(vector of dimension 6)

3	1	4	1
5	9	2	6
5	3	5	8
9	7	9	3
2	3	8	4
6	2	6	4

tensor of dimensions [6,4]
(matrix 6 by 4)

tensor of dimensions [4,4,2]

Logo Detection Using PyTorch

Project Pipeline

- 1. Get the data.**
- 2. Prepare data for network.**
- 3. Create network.**
- 4. Train network.**
- 5. Evaluate.**

1. Get the Data

- Flickrlogos-32
- 32 Different logo brands from Flickr

Adidas, Aldi, Apple, Becks, BMW, Carlsberg, Chimay, Coca-Cola, Corona, DHL, Erdinger, Esso, Fedex, Ferrari, Ford, Foster's, Google, Guinness, Heineken, HP, Milka, Nvidia, Paulaner, Pepsi, Ritter Sport, Shell, Singha, Starbucks, Stella Artois, Texaco, Tsingtao and UPS.

- Training logo set 320 images
- Validation logo set 960 images
- Test set 3,960 images
- No-logo set 3,000 images



<http://www.multimedia-computing.de/flickrlogos> (see download on page)

1. Get the Data

- Load FlickrLogos-32_dataset_v2.zip and unzip

```
import urllib.request
import os
import zipfile
import shutil
from pathlib import Path

FLICKLOGOS_URL = '.../FlickrLogos-32_dataset_v2.zip'
SOURCE_DIR = Path('FlickrLogos-v2')
DATA_DIR = Path('data')

def load_datasets(url, dst_dir):
    zip_file = url.split(sep='/')[-1]
    if not dst_dir.is_dir():
        if not zip_file.is_file():
            urllib.request.urlretrieve(url, zip_file)
        with zipfile.ZipFile(zip_file) as zip_ref:
            zip_ref.extractall()

load_datasets(FLICKLOGOS_URL, SOURCE_DIR)
```

2. Prepare Data for Network

- List image paths from text files and add half of no-logo paths to train and val paths

```
def list_image_paths(txt_relpath):  
    with open(txt_relpath) as f:  
        image_paths = f.read().splitlines()  
    return image_paths  
  
train_logo_relpaths = list_image_paths(SOURCE_DIR / 'trainset.relpaths.txt')  
val_logo_relpaths = list_image_paths(SOURCE_DIR / 'valset-logosonly.relpaths.txt')  
val_nologo_relpaths = list_image_paths(SOURCE_DIR / 'valset-nologos.relpaths.txt')  
test_relpaths = list_image_paths(SOURCE_DIR / 'testset.relpaths.txt')  
  
train_relpaths = train_logo_relpaths \  
    + val_nologo_relpaths[:int(len(val_nologo_relpaths) / 2)]  
val_relpaths = val_logo_relpaths \  
    + val_nologo_relpaths[int(len(val_nologo_relpaths) / 2):]
```

2. Prepare Data for Network

../adidas/2325670.jpg



../Dataset/Class/image.jpg
../train/adidas/2354545.jpg
../val/adidas/5553232.jpg
../test/adidas/7353256.jpg

```
SETS = ['train', 'val', 'test']
relpaths = [train_relpaths, val_relpaths, test_relpaths]
dataset_paths = dict(zip(SETS, relpaths))

def prepare_datasets(src_dir, dst_dir, keep_source=True):
    for dataset, paths in dataset_paths.items():
        num_files = 0
        for path in paths:
            num_files += 1
            src = src_dir / path
            dst = dst_dir / (path.replace('classes/jpg', dataset))
            dst.parent.mkdir(parents=True, exist_ok=True)
            shutil.copy2(src, dst)
        print(dataset, 'dataset:', str(num_files))
    if not keep_source: shutil.rmtree(src_dir)

prepare_datasets(SOURCE_DIR, DATA_DIR)
```


2. Prepare Data for Network

- Create data_transforms, datasets, and dataloaders.

```
import torch
import torchvision
from torchvision.transforms import transforms
from torch.utils.data import DataLoader

train_mean = np.array([0.44943, 0.4331, 0.40244])
train_std = np.array([0.29053, 0.28417, 0.30194])

data_transforms = transforms.Compose([
    transforms.Resize((224, 224)),
    transforms.ToTensor(),
    transforms.Normalize(train_mean, train_std)
])

datasets = {i : torchvision.datasets.ImageFolder(DATA_DIR / i, data_transforms)
            for i in SETS}

bz = 32

dataloaders = {i : DataLoader(datasets[i], batch_size=bz,
                              shuffle=(i == 'train'), num_workers=4) for i in SETS}
```

2. Prepare Data for Network

- Visualize datasets


```
import numpy as np
import matplotlib.pyplot as plt

def imshow(img):
    npimg = img.numpy().transpose((1, 2, 0))
    npimg = npimg * train_std + train_std # denorm
    npimg = np.clip(npimg, 0, 1)
    plt.imshow(npimg)

imgs, labels = next(iter(dataloaders['train']))


img = torchvision.utils.make_grid(imgs[:4])
classes = datasets['train'].classes
print(', '.join(classes[i] for i in labels[:4]))
imshow(img)
```

guinness, no-logo, carlsberg, no-logo



3. Create Network

Imitating LeNet-5



3. Create Network

- Create network by subclass torch.nn.Module

```
import torch.nn as nn
import torch.nn.functional as F

class CNN(nn.Module):
    def __init__(self):
        super(CNN, self).__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 53 * 53, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 33)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 16 * 53 * 53)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x
```

3. Create Network

- cnn

```
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
cnn = CNN()
cnn = cnn.to(device)
```

```
Out[51]: CNN(
  (conv1): Conv2d(3, 6, kernel_size=(5, 5), stride=(1, 1))
  (pool): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1, ceil_mode=False)
  (conv2): Conv2d(6, 16, kernel_size=(5, 5), stride=(1, 1))
  (fc1): Linear(in_features=44944, out_features=120, bias=True)
  (fc2): Linear(in_features=120, out_features=84, bias=True)
  (fc3): Linear(in_features=84, out_features=33, bias=True)
)
```

4. Train Network

- Define loss function and optimizer

```
import torch.optim as optim  
criterion = nn.CrossEntropyLoss()  
optimizer = optim.SGD(cnn.parameters(), lr=0.001, momentum=0.9)
```

4. Train Network

- Create train and validate function

```
def train_val(model, criterion, optimizer, num_epochs=25):  
    since = time.time()  
  
    best_model_wts = copy.deepcopy(model.state_dict())  
    best_acc = 0.0  
  
    for epoch in range(num_epochs):  
        print('Epoc {}/{}'.format(epoch, num_epochs - 1))  
        print('-' * 10)  
        for phase in ['train', 'val']:  
            if phase == 'train':  
                model.train()  
            else:  
                model.eval()  
            running_loss = 0.0  
            running_corrects = 0  
            for inputs, labels in dataloaders[phase]:  
                inputs, labels = inputs.to(device), labels.to(device)
```

4. Train Network

- Create train and validate function

```
def train_val(model, criterion, optimizer, num_epochs=25):  
    ...  
    for inputs, labels in dataloaders[phase]:  
        inputs, labels = inputs.to(device), labels.to(device)  
        optimizer.zero_grad()  
        with torch.set_grad_enabled(phase == 'train'):  
            # Forward pass  
            outputs = model(inputs)  
            _, preds = torch.max(outputs, 1)  
            # Compute loss  
            loss = criterion(outputs, labels)  
            # Compute gradients and update parameters if train  
            if phase == 'train':  
                loss.backward()  
                optimizer.step()
```


4. Train Network

- Create train and validate function

```
def train_val(model, criterion, optimizer, num_epochs=25):  
    ...  
    time_elapsed = time.time() - since  
    print('Training complete in {:.0f}m {:.0f}s'.format(  
        time_elapsed // 60, time_elapsed % 60))  
    print('Best Accuracy: {:.2f} %'.format(best_acc * 100))  
  
    model.load_state_dict(best_model_wts)  
  
    return model
```

4. Train Network

- Train the network

```
model_cnn = train_val(cnn, criterion, optimizer)
```

```
Epoc 0/24
-----
Train Loss: 1.8010 Acc.: 75.71 %
Val Loss: 2.2365 Acc.: 60.98 %

Epoc 1/24
-----
Train Loss: 1.1586 Acc.: 82.42 %
Val Loss: 2.0667 Acc.: 60.98 %
...

Epoc 20/24
-----
Train Loss: 0.5419 Acc.: 86.98 %
Val Loss: 3.0452 Acc.: 61.87 %
...

Epoc 24/24
-----
Train Loss: 0.3455 Acc.: 92.14 %
Val Loss: 2.6079 Acc.: 57.72 %

Training complete in 11m 4s
Best Accuracy: 61.87 %
```

5. Evaluate

- Create test function to evaluate network on test set

```
def test(model):
    model.eval()
    running_corrects = 0
    with torch.no_grad():
        for inputs, labels in dataloaders['test']:
            inputs, labels = inputs.to(device), labels.to(device)
            outputs = model(inputs)
            _, preds = torch.max(outputs, 1)

            running_corrects += torch.sum(preds == labels).item()

    test_acc = running_corrects / len(datasets['test'])
    print('Test Acc.: {:.2f} %'.format(test_acc * 100))

test(model_cnn)
```

Test Acc.: 75.91 %

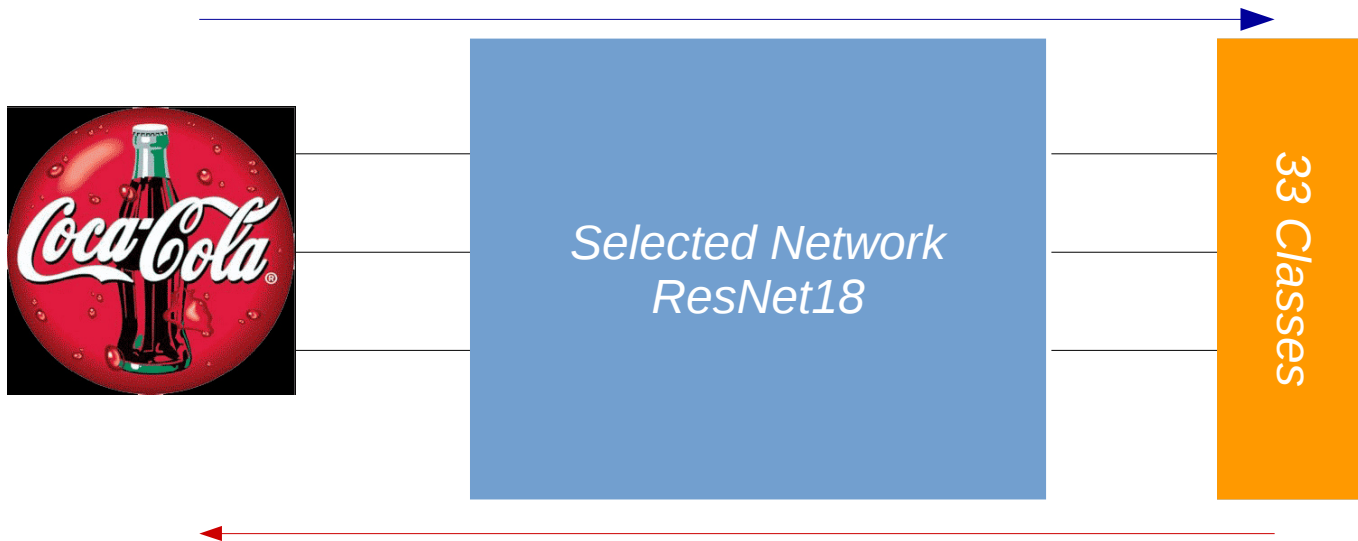
Transfer Learning

What's Transfer Learning?

“Transfer learning is a machine learning technique where a model trained on one task is re-purposed on a second related task.”

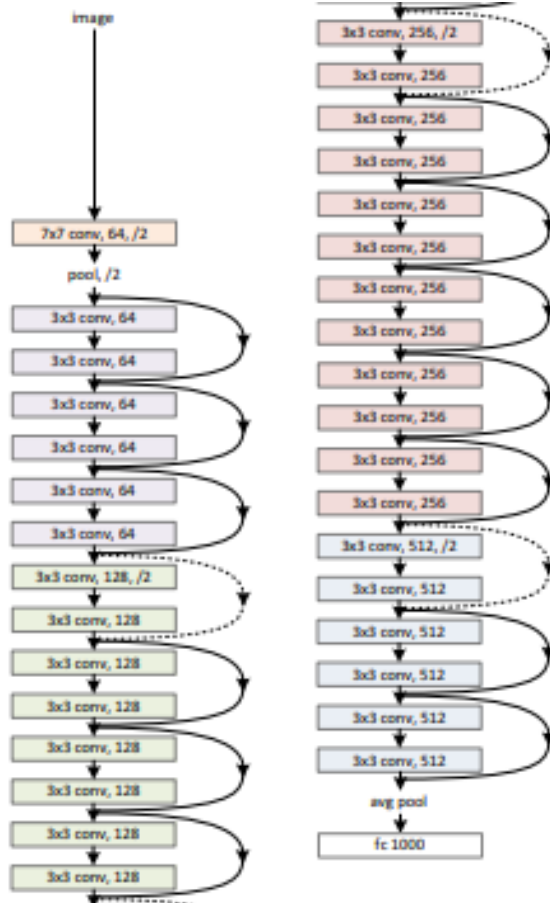
How?

1. Select network.
2. Match input format.
3. Replace output layer.
4. Retrain network



ResNet

34-layer residual



layer name	output size	18-layer	34-layer	50-layer	101-layer	152-layer
conv1	112×112	7×7, 64, stride 2				
		3×3 max pool, stride 2				
conv2_x	56×56	$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 64 \\ 3 \times 3, 64 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 64 \\ 3 \times 3, 64 \\ 1 \times 1, 256 \end{bmatrix} \times 3$
conv3_x	28×28	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 128 \\ 3 \times 3, 128 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 4$	$\begin{bmatrix} 1 \times 1, 128 \\ 3 \times 3, 128 \\ 1 \times 1, 512 \end{bmatrix} \times 8$
conv4_x	14×14	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 256 \\ 3 \times 3, 256 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 6$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 23$	$\begin{bmatrix} 1 \times 1, 256 \\ 3 \times 3, 256 \\ 1 \times 1, 1024 \end{bmatrix} \times 36$
conv5_x	7×7	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 2$	$\begin{bmatrix} 3 \times 3, 512 \\ 3 \times 3, 512 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$	$\begin{bmatrix} 1 \times 1, 512 \\ 3 \times 3, 512 \\ 1 \times 1, 2048 \end{bmatrix} \times 3$
	1×1	average pool, 1000-d fc, softmax				
FLOPs		1.8×10^9	3.6×10^9	3.8×10^9	7.6×10^9	11.3×10^9

Match Input Format

- Create matched transforms

```
train_mean = np.array([0.485, 0.456, 0.406])  
train_std = np.array([0.229, 0.224, 0.225])
```

```
data_transforms = transforms.Compose([  
    transforms.Resize((224, 224)),  
    transforms.ToTensor(),  
    transforms.Normalize(train_mean, train_std)])
```

```
datasets = {i : torchvision.datasets.ImageFolder(DATA_DIR / i, data_transforms)  
            for i in SETS}
```

```
dataloaders = {i : DataLoader(datasets[i], batch_size=bz,  
                              shuffle=(i == 'train'), num_workers=4) for i in SETS}
```


Load Pretrained Network

- Load pretrained model. Define output layer (head).

```
model_ft = torchvision.models.resnet18(pretrained=True)
model_ft
```

```
Out[19]: ResNet(
  (conv1): Conv2d(3, 64, kernel_size=(7, 7), stride=(2, 2), padding=(3, 3), bias=False)
  (bn1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)
  (relu): ReLU(inplace)
  (maxpool): MaxPool2d(kernel_size=3, stride=2, padding=1, dilation=1, ceil_mode=False)
  (layer1): Sequential(
    ...
  )
  )
  (avgpool): AvgPool2d(kernel_size=7, stride=1, padding=0)
  (fc): Linear(in_features=512, out_features=1000, bias=True)
)
```

Replace Output Layer

- Replace output layer.

```
num_fts = model_ft.fc.in_features          # 512
model_ft.fc = nn.Linear(num_fts, 33)
model_ft = model_ft.to(device)
criterion = nn.CrossEntropyLoss()
lr = 0.001
optimizer_ft = optim.SGD(model_ft.parameters(), lr=lr, momentum=0.9)
```

Train Network

- Fine tune the network

```
model_ft = train_val(model_ft, criterion, optimizer_ft)
```

```
Epoc 0/24
-----
Train Loss: 1.0930 Acc.: 78.35 %
Val Loss: 1.6996 Acc.: 61.38 %

Epoc 1/24
-----
Train Loss: 0.6551 Acc.: 83.30 %
Val Loss: 1.6968 Acc.: 61.87 %
...

Epoc 22/24
-----
Train Loss: 0.0142 Acc.: 100.00 %
Val Loss: 1.1601 Acc.: 75.69 %
...

Epoc 24/24
-----
Train Loss: 0.0132 Acc.: 99.95 %
Val Loss: 1.1957 Acc.: 75.65 %

Training complete in 13m 9s
Best Accuracy: 75.69 %
```

Evaluate

- Evaluate network on test set

```
test(model_ft)
```

Test Acc.: 85.58 %

Thank you

All materials:

<https://github.com/nithiroj/pycon-thailand-2018>

NITHIROJ TRIPATARASIT (Lek)

Email: nithiroj@gmail.com

Twitter: [@nithiroj](https://twitter.com/nithiroj)

URL: medium.com/@nithiroj

GitHub: github.com/nithiroj

LinkedIn: [nithiroj-tripatarasit](https://www.linkedin.com/company/nithiroj-tripatarasit)